

CS 4350: Fundamentals of Software Engineering

Lecture 5: Concurrency Patterns in Typescript

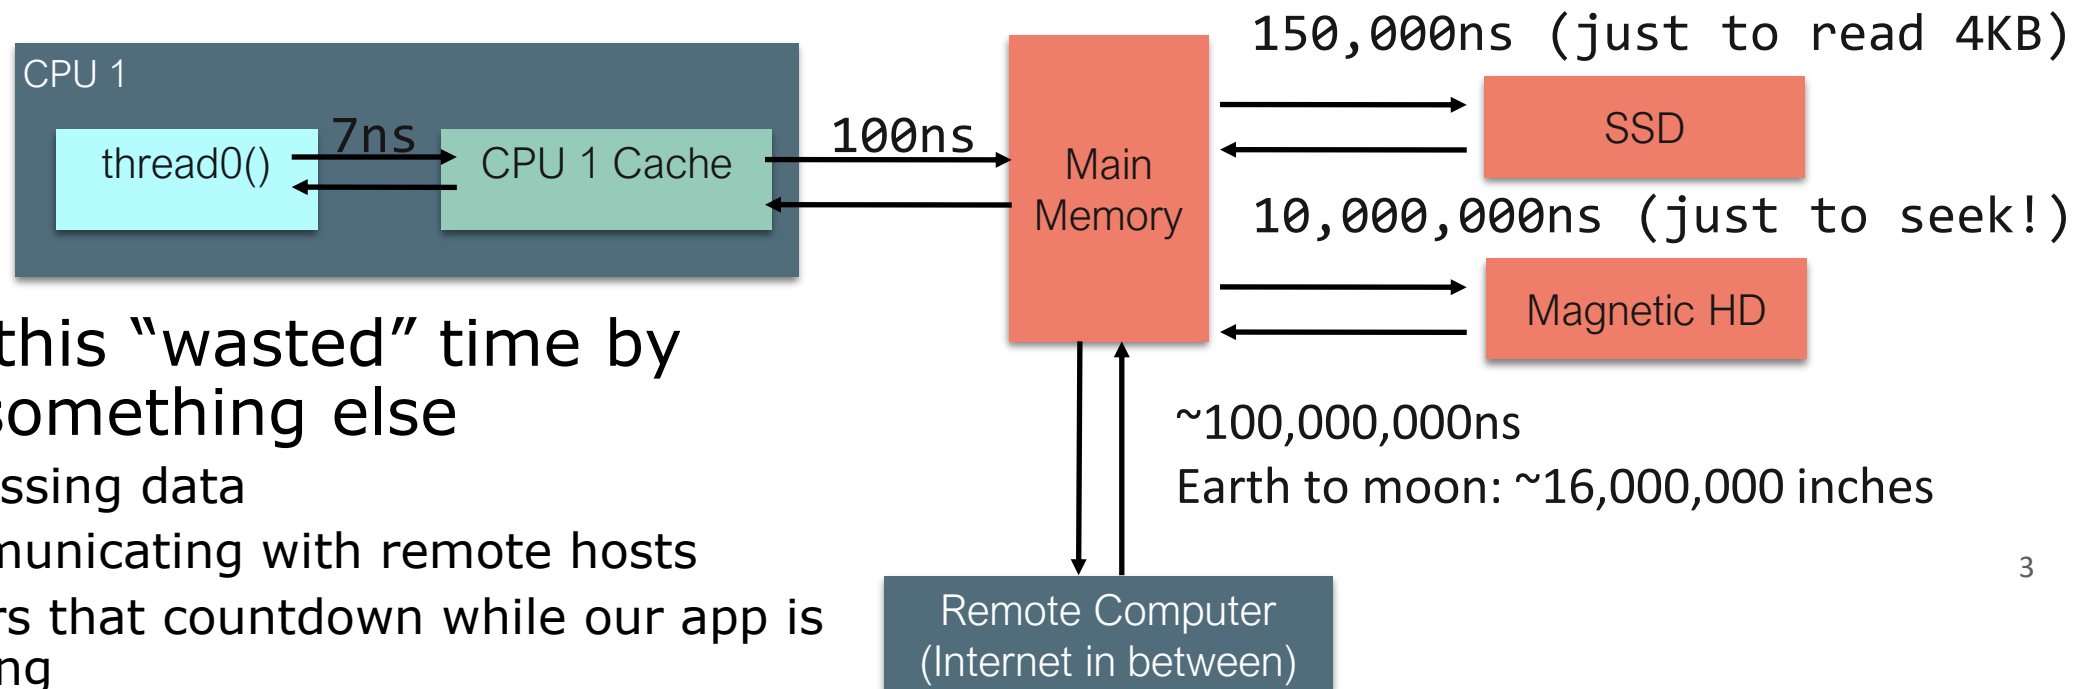
Jonathan Bell, Adeel Bhutta, Mitch Wand
Khoury College of Computer Sciences

Learning Goals for this Lesson

- At the end of this lesson, you should be prepared to:
 - Explain how to achieve concurrency through asynchronous operations and `Promise.all` in TypeScript.
 - Write asynchronous and concurrent code in TypeScript using `async/await` and `Promise.all`.

Masking Latency with Concurrency

- Consider: a 1Ghz CPU executes an instruction every 1 ns
- Almost anything else takes forever (approximately)



- Utilize this “wasted” time by doing something else
 - Processing data
 - Communicating with remote hosts
 - Timers that countdown while our app is running
 - Waiting for users to provide input

Pre-emptive Multiprocessing

- OS manages multiprocessing with multiple threads of execution
- Processes may be interrupted at unpredictable times
- Interprocess communication by shared memory
- Data races abound
- Really, really hard to get right: need critical sections, semaphores, monitors (all that stuff you learned about in op. sys.)

An alternative model: cooperative multiprocessing

- OS manages multiprocessing with multiple threads of execution
- Each thread decides when it should yield to let other threads execute
- Typically via a **yield** or **await** operation

JavaScript/TypeScript implements Cooperative Multiprocessing Using “run-to-completion” semantics

- JS has primitives that allow one computation to start another computation that runs concurrently with the first.
- These are almost always IO operations.
- However, the original computation **always runs to completion.**


Run-to-completion semantics

- A computation runs continuously until it is either suspended or completed.
 - This means that only one of your computations is running at any time (in addition to whatever asynchronous IO is running)
- A computation is suspended when it hits an `'await'`. The runtime system (node.js, for us) chooses what to do next. (In addition to whatever asynchronous IO it may be doing).

Defining a concurrent computation

```
async function makeOneGetRequest(requestNumber:number) {  
  const response = await axios.get('https://rest-example.covey.town');  
  console.log(`For request ${requestNumber}, server replied: `,  
response.data);  
}
```

- An **async function** is a function that creates a concurrent computation.
- Calling the function will tell the operating system to start the computation.
- TS vocabulary: This computation is called a **promise**



This is the address of a server that returns the number of calls that have been made to this server.

One concurrent computation can wait for the result of another one.


```
async function makeOneGetRequest(requestNumber:number) {  
    const response = await axios.get('https://rest-example.covey.town');  
    console.log(`For request ${requestNumber}, server replied: `,  
response.data);  
}
```

- Axios.get is also an async function, so it returns a promise (let's call it **p**)
- The **await** suspends the current computation until the promise **p** returns.
- While the current computation is suspended, other computations (including **p**) can run.

Example:

```
async function makeThreeSimpleRequests() {  
  makeOneGetRequest(1);  
  makeOneGetRequest(2);  
  makeOneGetRequest(3);  
  console.log("Three requests made")  
}
```

```
makeThreeSimpleRequests()
```



```
$ npx ts-node example2.ts
```

```
Three requests made
```

```
For request 2, server replied: This is GET number 280 on the current server
```


```
For request 3, server replied: This is GET number 281 on the current server
```

```
For request 1, server replied: This is GET number 282 on the current server
```

Awaiting a promise prevents your method from continuing

```
async function makeThreeSerialRequests(): Promise<void> {  
    await makeOneGetRequest(1);  
    await makeOneGetRequest(2);  
    await makeOneGetRequest(3);  
    console.log('Heard back from all of the requests')  
}
```

```
makeThreeSerialRequests();
```



```
For request 1, server replied: This is GET number 37 on the current server  
For request 2, server replied: This is GET number 38 on the current server  
For request 3, server replied: This is GET number 39 on the current server  
Heard back from all of the requests  
Elapsed time: 364.0822000205517 milliseconds
```


Promise.all starts several promises concurrently

```
async function makeThreeConcurrentRequests(): Promise<void> {  
    await Promise.all([  
        makeOneGetRequest(1),  
        makeOneGetRequest(2),  
        makeOneGetRequest(3)  
    ])  
    console.log('Heard back from all of the requests')  
}
```

- **Promise.all** takes a list of promises and runs them all concurrently.
- It finishes when all the promises have finished.

Promise.all allows for concurrency

```
async function makeThreeConcurrentRequests(): Promise<void> {  
    await Promise.all([  
        makeOneGetRequest(1),  
        makeOneGetRequest(2),  
        makeOneGetRequest(3)  
    ])  
    console.log('Heard back from all of the requests')  
}  
  
makeThreeConcurrentRequests();
```



```
For request 2, server replied: This is GET number 58 on the current server  
For request 1, server replied: This is GET number 59 on the current server  
For request 3, server replied: This is GET number 60 on the current server  
Heard back from all of the requests  
Elapsed time: 203.7674999833107 milliseconds
```

Visualizing Promise.all (1)

Sequential version: ~400msec

```
async function makeThreeSerialRequests():  
Promise<void> {  
    await makeOneGetRequest(1);  
    await makeOneGetRequest(2);  
    await makeOneGetRequest(3);  
    console.log('Heard back from all of the  
requests')  
}
```

“Don’t make another request
until you got the last response
back”

Concurrent version: ~126msec

```
async function makeThreeConcurrentRequests():  
Promise<void> {  
    await Promise.all([  
        makeOneGetRequest(1),  
        makeOneGetRequest(2),  
        makeOneGetRequest(3)  
    ])  
    console.log('Heard back from all of the requests')  
}
```

“Make all of the requests at
the same time, then wait for
all of the responses”

Visualizing Promise.all (2)


Sequential version: ~400msec

```
async function makeThreeSerialRequests():  
  Promise<void> {  
    await makeOneGetRequest(1);  
    await makeOneGetRequest(2);  
    await makeOneGetRequest(3);  
    console.log('Heard back from all of the  
requests')  
  }
```

Concurrent version: ~126msec

```
async function makeThreeConcurrentRequests():  
  Promise<void> {  
    await Promise.all([  
      makeOneGetRequest(1),  
      makeOneGetRequest(2),  
      makeOneGetRequest(3)  
    ])  
    console.log('Heard back from all of the requests')  
  }
```

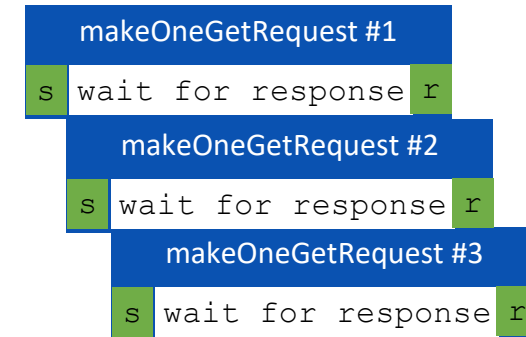
Time →



s send

r receive

Time →



Patterns for Concurrent Code: Example: Using a Web Service

Here is a web service
we'd like to talk to.

```
POST /transcripts
-- adds a new student to the database,
-- returns an ID for this student.
-- requires a body parameter 'name'
-- Multiple students may have the same name.
GET /transcripts/:ID
-- returns transcript for student with given ID. Fails if no such student
DELETE /transcripts/:ID
-- deletes transcript for student with the given ID, fails if no such student
POST /transcripts/:studentID/:courseNumber
-- adds an entry in this student's transcript with given name and course.
-- Requires a body parameter 'grade'
-- Fails if there is already an entry for this course in the student's transcript
GET /transcripts/:studentID/:courseNumber
-- returns the student's grade in the specified course.
-- Fails if student or course is missing.
GET /studentids?name=string
-- returns list of IDs for student with the given name
```


An Example Task Using the Transcript Server

- Given an array of StudentIDs:
 - Request each student's transcript, and save it to disk so that we have a copy
 - Once all of the pages are downloaded and saved, print out the total size of all of the files that were saved

Generating a promise for a student

```
async function promiseForTranscript(studentID: number) {  
  const response = await axios.get(`https://rest-example.covey.town/transcripts/${studentID}`)
```

}
Here is something we plan
to do later

The promise is to call axios
and wait for the result.

Generating a promise for a student (cont'd)

```
async function promiseForTranscript(studentID: number) {  
  const response = await axios.get(`https://rest-example.covey.town/transcripts/${studentID}`)  
  await fsPromises.writeFile(`transcript-${response.data.student.studentID}.json`,  
    JSON.stringify(response.data))  
}
```

After we get the response, make a new promise: this time to write the result to a file. Then wait for that to finish.

When the file-writing promise is fulfilled, then the whole original promise is fulfilled.

Now, actually generate all the promises

```
async function runClientAsync(studentIDs:number[]) {  
  console.log('Making requests for ${studentIDs}');
```

```
  async function promiseForTranscript(studentID: number) { .. }
```

```
  const promisesForTranscripts = studentIDs.map(promiseForTranscript)  
  console.log('Requests sent!');
```

map applies the function specified to each element in the array and returns a new array containing the result of each of those functions

```
}
```

Wait for all the promises to resolve

```
async function runClientAsync(studentIDs:number[]) {  
  console.log('Making requests for ${studentIDs}');  
  
  async function promiseForTranscript(studentID: number) { .. }  
  
  const promisesForTranscripts = studentIDs.map(promiseForTranscript)  
  console.log('Requests sent!');  
  await Promise.all(promisesForTranscripts);  
  
}
```

Asynchronously stat all the files

```
async function runClientAsync(studentIDs:number[]) {  
  console.log('Making requests for ${studentIDs}');  
  
  async function promiseForTranscript(studentID: number) { .. }  
  
  const promisesForTranscripts = studentIDs.map(promiseForTranscript)  
  console.log('Requests sent!');  
  await Promise.all(promisesForTranscripts);  
  const stats = await Promise.all(studentIDs.map(studentID => fsPromises.stat(`transcript-  
${studentID}.json`)));  
  
}
```

..then total the sizes

```
async function runClientAsync(studentIDs:number[]) {  
  console.log('Making requests for ${studentIDs}');  
  
  async function promiseForTranscript(studentID: number) { .. }  
  
  const promisesForTranscripts = studentIDs.map(promiseForTranscript)  
  console.log('Requests sent!');  
  await Promise.all(promisesForTranscripts);  
  const stats = await Promise.all(studentIDs.map(studentID => fsPromises.stat(`transcript-${studentID}.json`)));  
  const totalSize = stats.reduce((runningTotal, val) => runningTotal + val.size, 0);  
  console.log(`Finished calculating size: ${totalSize}`);  
  console.log('Done');  
}
```

'reduce' is what you called 'foldl' back in Fundies 1.

Leverage Concurrency When Possible

• Where you place awaits can make a big difference! ✓

For each student: make an async handler to fetch their transcript and save it

The code we've seen on past slides:

```
async function runClientAsync() {
  console.log('Making a requests');
  const studentIDs = [1, 2, 3, 4];
  const promisesForTranscripts = studentIDs.map(
    async (studentID) => {
      const response = await axios.get(`https://rest-example.covey.town/transcripts/${studentID}`)
      await fsPromises.writeFile(`transcript-${response.data.student.studentID}.json`, JSON.stringify(response.data))
    });
  console.log('Requests sent!');
  await Promise.all(promisesForTranscripts);
  const stats = await Promise.all(studentIDs.map(studentID => fsPromises.stat(`transcript-${studentID}.json`)));
  const totalSize = stats.reduce((runningTotal, val) => runningTotal + val.size, 0);
  console.log('Finished calculating size: ${totalSize}');
}
```

Running time:
1.5 sec

For each student: wait to fetch their transcript, then wait to write it, then go on to the next student

This accomplishes the same function, but without concurrency:

```
async function runClientAsyncSerially() {
  console.log('Making a requests');
  const studentIDs = [1, 2, 3, 4];
  for(let studentID of studentIDs){
    const response = await axios.get(`https://rest-example.covey.town/transcripts/${studentID}`);
    await fsPromises.writeFile(`transcript-${response.data.student.studentID}.json`, JSON.stringify(response.data))
  }
  let totalSize = 0;
  for(let studentID of studentIDs){
    const stats = await fsPromises.stat(`transcript-${studentID}.json`);
    totalSize += stats.size;
  }
  console.log('Finished calculating size: ${totalSize}');
}
```

Running time:
2.2 sec

This is what we mean by “your code can become synchronous”

Async/Await Programming Activity

- Your task is to write a new `async` function, `importGrades`, which takes in input of the type `ImportTranscript[]`.
- `importGrades` should create a student record for each `ImportTranscript`, and then post the grades for each of those students.
- After posting the grades, it should fetch the transcripts for each student and return an array of transcripts.

Download the activity (includes instructions in README.md):
Linked from course webpage for week 4, or at <https://bit.ly/34GbcN6>

Learning Goals for this Lesson

- At the end of this lesson, you should be prepared to:
 - Explain how to achieve concurrency through asynchronous operations and `Promise.all` in TypeScript.
 - Write asynchronous and concurrent code in TypeScript using `async/await` and `Promise.all`.

Learning Goals for this Lesson (expanded)

- At the end of this lesson, you should be prepared to:
 - Explain how to achieve concurrency through asynchronous operations and `Promise.all` in TypeScript.
 - Write asynchronous and concurrent code in TypeScript using `async/await` and `Promise.all`.
 - Write asynchronous code using promises and `.then()`.
 - Explain the difference between JS run-to-completion semantics and interrupt-based semantics.

Additional Topics

General Rules for Writing Asynchronous Code

- Don't perform long-running computations or synchronous IO
- Leverage concurrency when possible
 - Remember that events are processed in the order they are **received**
 - But events may arrive in unexpected order!
- Always check for errors (try/catch for async/await, ".catch" for promises)

Async functions use Promises Under the Hood


Promises Enforce Ordering Through “Then”

```
1. console.log('Making requests');
2. axios.get('https://rest-example.covey.town/')
   .then((response) =>{
       console.log('Heard back from server');
       console.log(response.data);
   });
3. axios.get('https://www.google.com/')
   .then((response) =>{
       console.log('Heard back from Google');
   });
4. axios.get('https://www.facebook.com/')
   .then((response) =>{
       console.log('Heard back from Facebook');
   });
5. console.log('Requests sent!');
```

- **axios.get** returns a promise.
- **p.then** mutates that promise so that the then block is run immediately after the original promise returns.
- The resulting promise isn't completed until the then block finishes.
- You can chain **.then's**, to get things that look like `p.then().then().then()`

Async/await code is compiled into promise/then code

```
async function
makeThreeSerialRequests() {
1.  console.log('Making first
request');
2.  await makeOneGetRequest();
3.  console.log('Making second
request');
4.  await makeOneGetRequest();
5.  console.log('Making third
request');
6.  await makeOneGetRequest();
7.  console.log('All done!');
}
makeThreeSerialRequests();
```



```
console.log('Making first request');
makeOneGetRequest().then(() => {
  console.log('Making second request');
  return makeOneGetRequest();
}).then(() => {
  console.log('Making third request');
  return makeOneGetRequest();
}).then(() => {
  console.log('All done!');
});
```


Syntax for Writing Asynchronous Code

- You can only call **await** from a function that is **async**
- You can only **await** on functions that return a **Promise**
- Beware: **await** makes your code synchronous (this is what we want it for)!
- Handle errors using try/catch instead of "catch" (common gotcha with promises)

```
async function makeOneGetRequest(): Promise<void> {  
    console.log("Making Request");  
    try {  
        const response = await axios.get("https://rest-  
example.covey.town");  
        console.log("Heard back from server");  
        console.log(response.data);  
    } catch (err) {  
        console.log('Uh oh!'); console.trace(err);  
    }  
}
```

```
function makeOneGetRequestNoAsync(): Promise<void> {  
    console.log("Making Request");  
    return axios.get("https://rest-  
example.covey.town").then((response) => {  
        console.log("Heard back from server");  
        console.log(response.data);  
    }).catch(err => {  
        console.log('Uh oh!');  
        console.trace(err);  
    });  
}
```

Data Races in TS vs. Java

Data Races in TS vs. Java

```
let x : number = 1
```

```
async function asyncDouble() {  
    // start an asynchronous computation and wait for the result  
    await makeOneGetRequest(1);  
    x = x * 2 // statement 1  
}
```

```
async function asyncIncrementTwice() {  
    // start an asynchronous computation and wait for the result  
    await makeOneGetRequest(2);  
    x = x + 1; // statement 2  
    x = x + 1; // statement 3  
}
```

```
async function run() {  
    await Promise.all([asyncDouble(), asyncIncrementTwice()])  
    console.log(x)  
}
```

Explanation

- In the JS run-to-completion semantics, statement 3 is guaranteed to run immediately after statement 2, so the only possible orders of execution are:
 - 1,2,3 (1 runs before 2 and 3, final value of x is 4)
 - 2,3,1 (2 and 3 run before 1, final value of x is 6)
- In an interrupt-based model, it is possible that statement 1 runs **BETWEEN** statement 2 and statement 3, yielding the order of execution
 - 2,1,3 (final value of x is 5).

Explanation (2)

- Notice that there is still a data race between statement 1 and statements 2 and 3;
- Run-to-completion semantics does not eliminate data races entirely, but it makes them much rarer.

Learning Goals for this Lesson (expanded)

- At the end of this lesson, you should be prepared to:
 - Explain how to achieve concurrency through asynchronous operations and `Promise.all` in TypeScript.
 - Write asynchronous and concurrent code in TypeScript using `async/await` and `Promise.all`.
 - Write asynchronous code using promises and `.then()`.
 - Explain the difference between JS run-to-completion semantics and interrupt-based semantics.